

CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

TIME AND ORDERING

Lecture D

LAMPORT TIMESTAMPS

ORDERING EVENTS IN A DISTRIBUTED SYSTEM

- To order events across processes, trying to sync clocks is one approach.
- What if we instead assigned timestamps to events that were not *absolute* time?
- As long as these timestamps obey *causality*, that would work.

If an event A causally happens before another event B, then $\text{timestamp}(A) < \text{timestamp}(B)$.

Humans use causality all the time.

E.g., I enter a house only after I unlock it.

E.g., you receive a letter only after I send it.


LOGICAL (OR LAMPORT) ORDERING

- Proposed by Leslie Lamport in the 1970s
- Used in almost all distributed systems since then
- Almost all cloud computing systems use some form of logical ordering of events

LOGICAL (OR LAMPORT) ORDERING(2)

- Define a logical relation *Happens-Before* among pairs of events
- *Happens-Before* denoted as \rightarrow
- Three rules
 1. On the same process: $a \rightarrow b$, if $time(a) < time(b)$ (using the local clock)
 2. If p1 sends m to p2: $send(m) \rightarrow receive(m)$
 3. (Transitivity) If $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$
- Creates a *partial order* among events
 - Not all events related to each other via \rightarrow

EXAMPLE




While P1 and P3 each have an event labeled E, these are different events as they occur at different processes

● *Instruction or step*


→ *Message*

HAPPENS-BEFORE



- $A \rightarrow B$
- $B \rightarrow F$
- $A \rightarrow F$

HAPPENS-BEFORE (2)



P3.

- $H \rightarrow G$
- $F \rightarrow J$
- $H \rightarrow J$
- $C \rightarrow J$

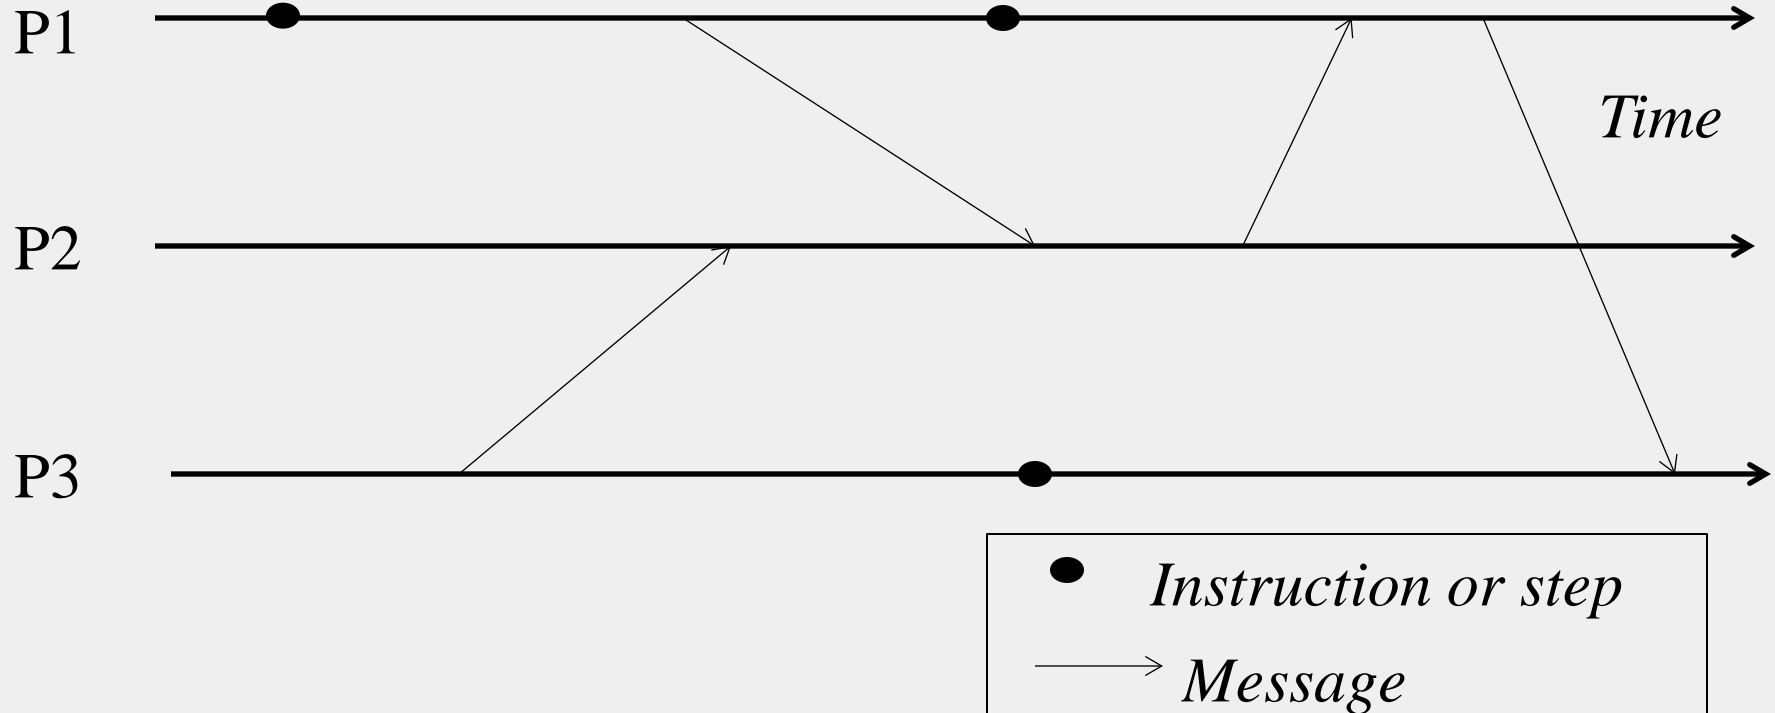
● *Instruction or step*

→ *Message*

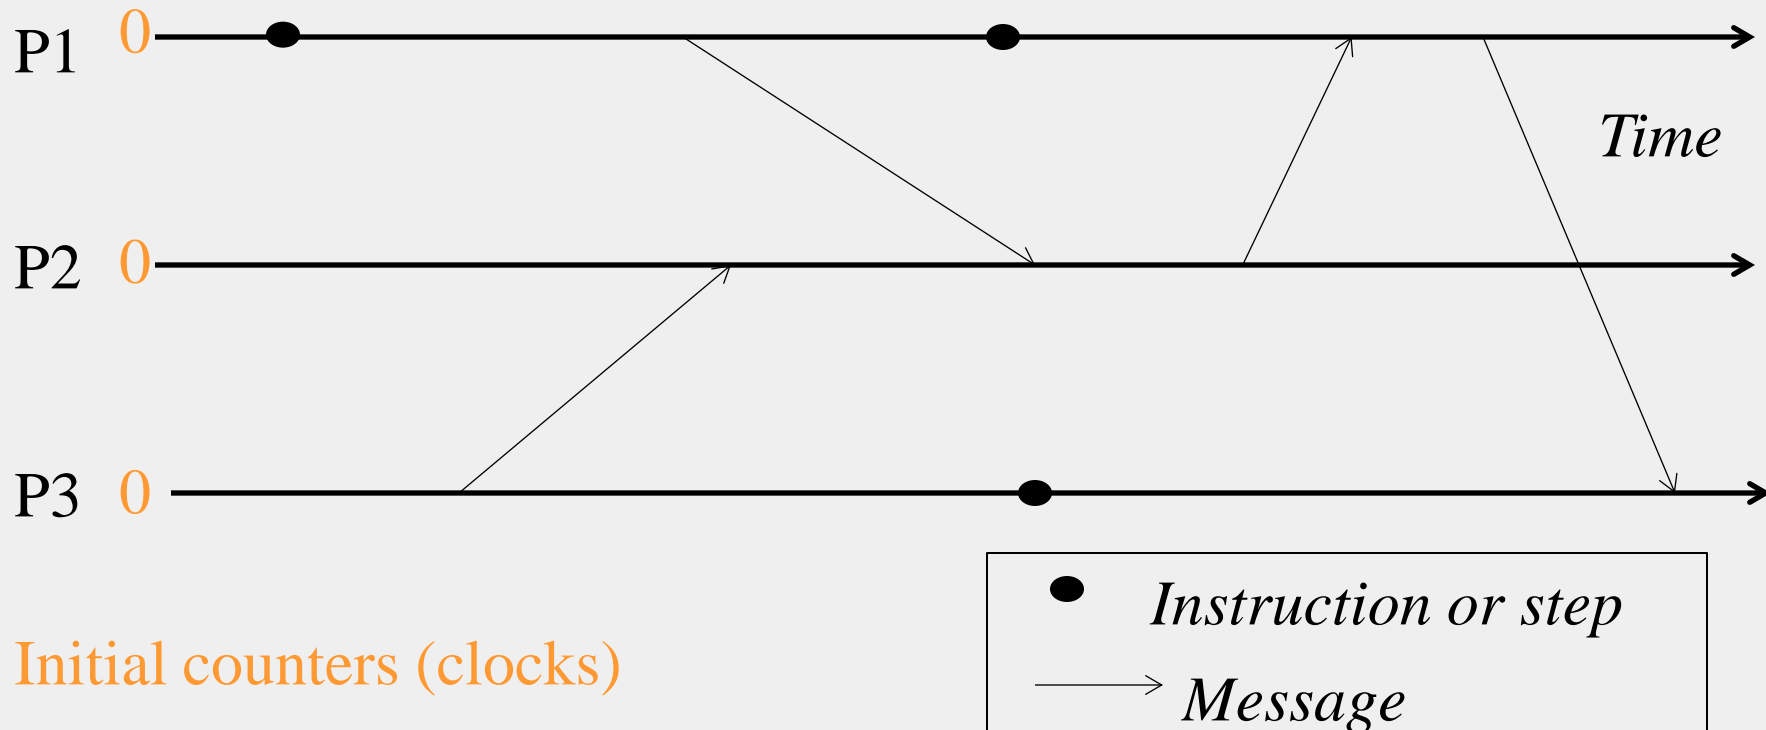
IN PRACTICE: LAMPORT TIMESTAMPS

- **Goal:** Assign logical (Lamport) timestamp to each event
- **Timestamps obey causality**
- **Rules**
 - Each process uses a local counter (clock) which is an integer
 - **Initial** value of counter is **zero**
 - A process **increments** its counter when a **send** or an **instruction** happens at it. The counter is assigned to the event as its timestamp.
 - A **send (message)** event carries its timestamp
 - For a **receive (message)** event the counter is updated by
$$\max(\text{local clock, message timestamp}) + 1$$

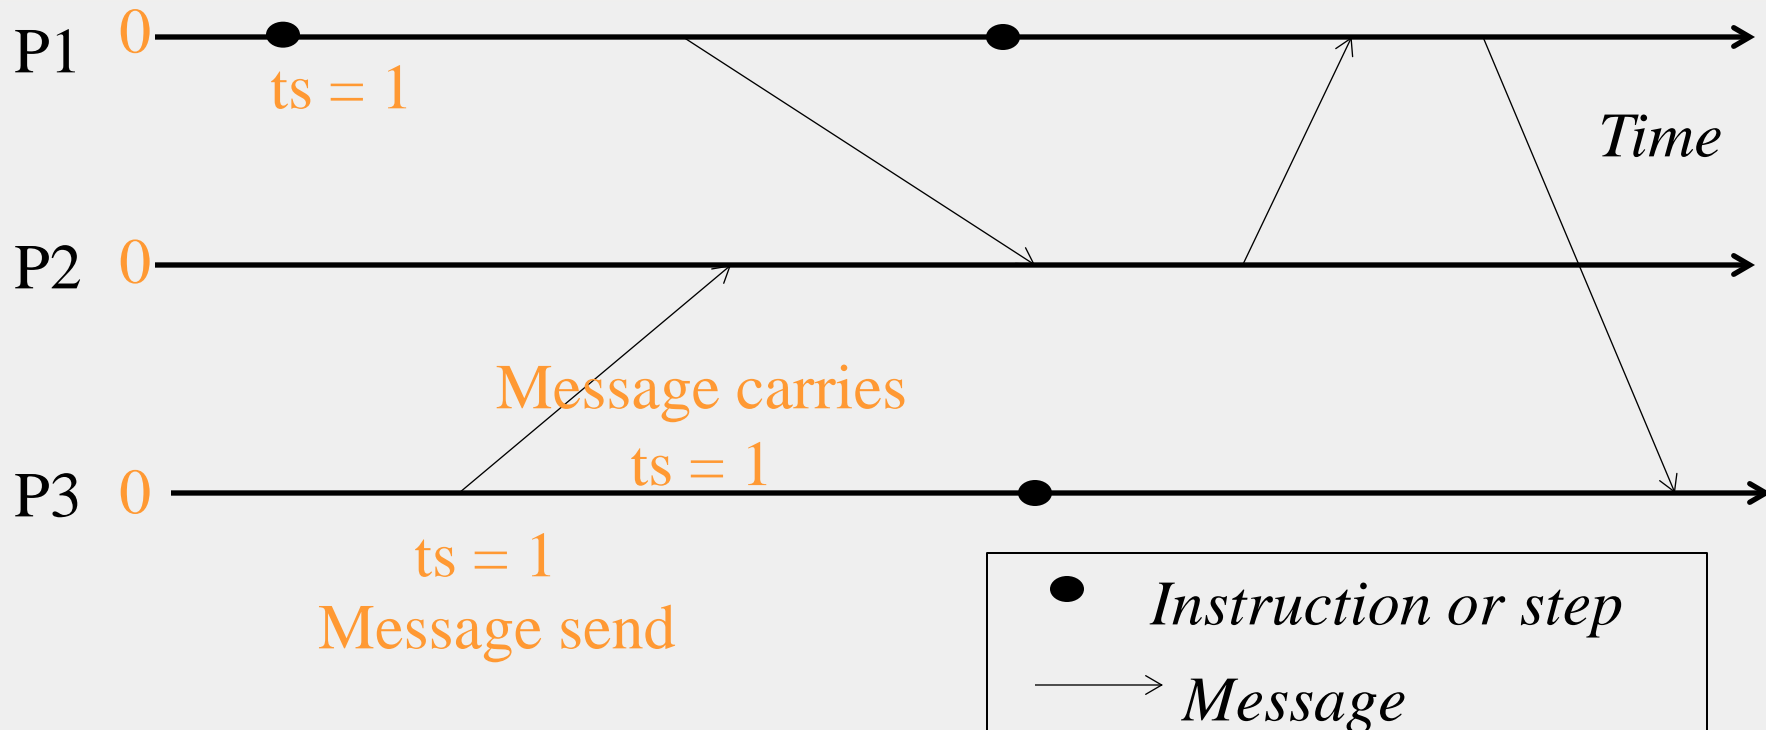
EXAMPLE



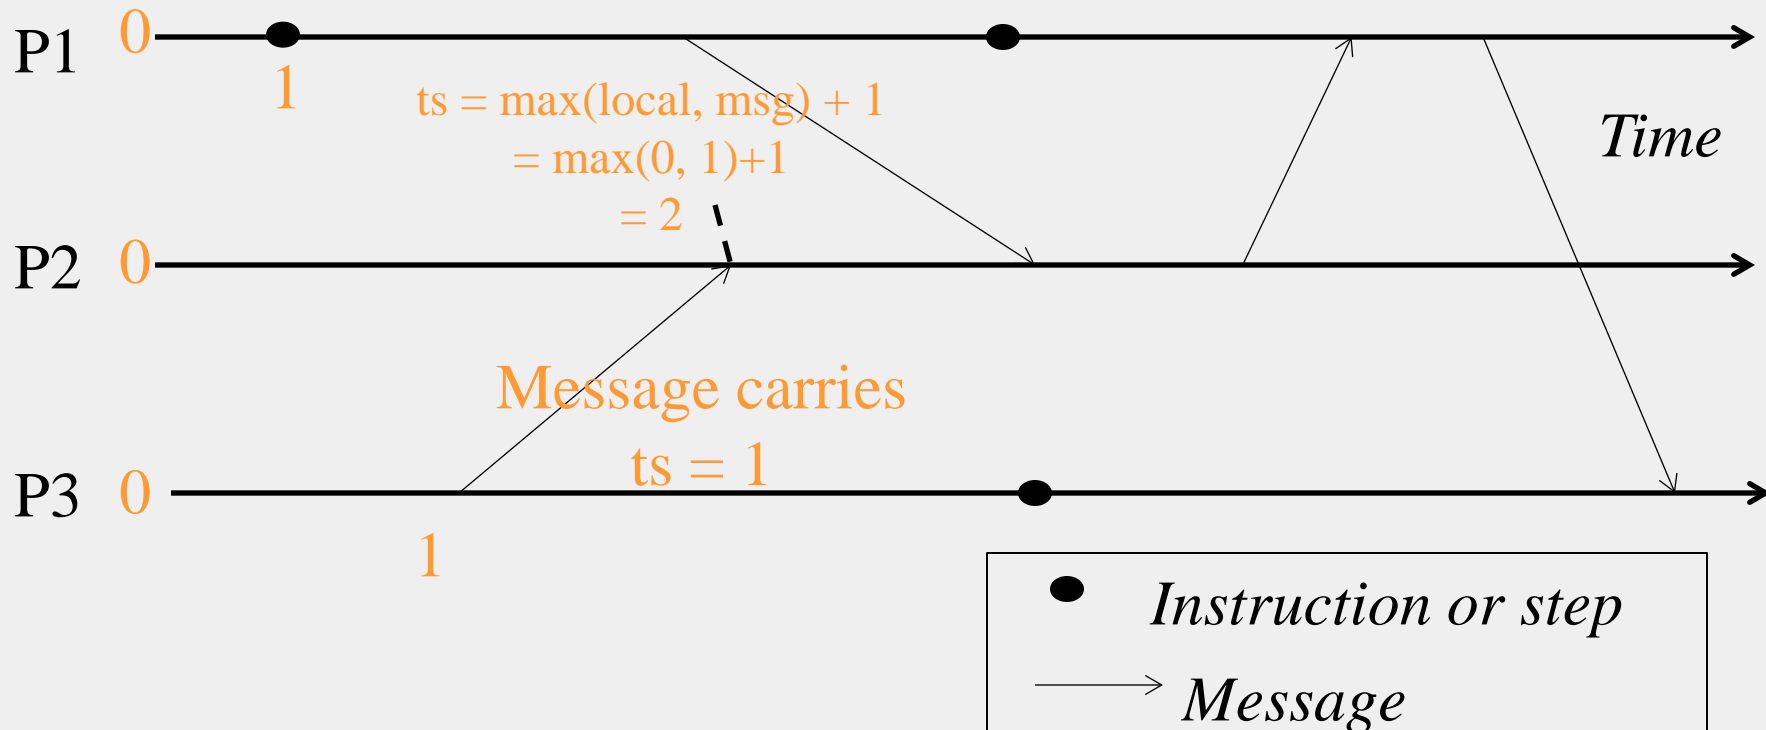
LAMPORT TIMESTAMPS



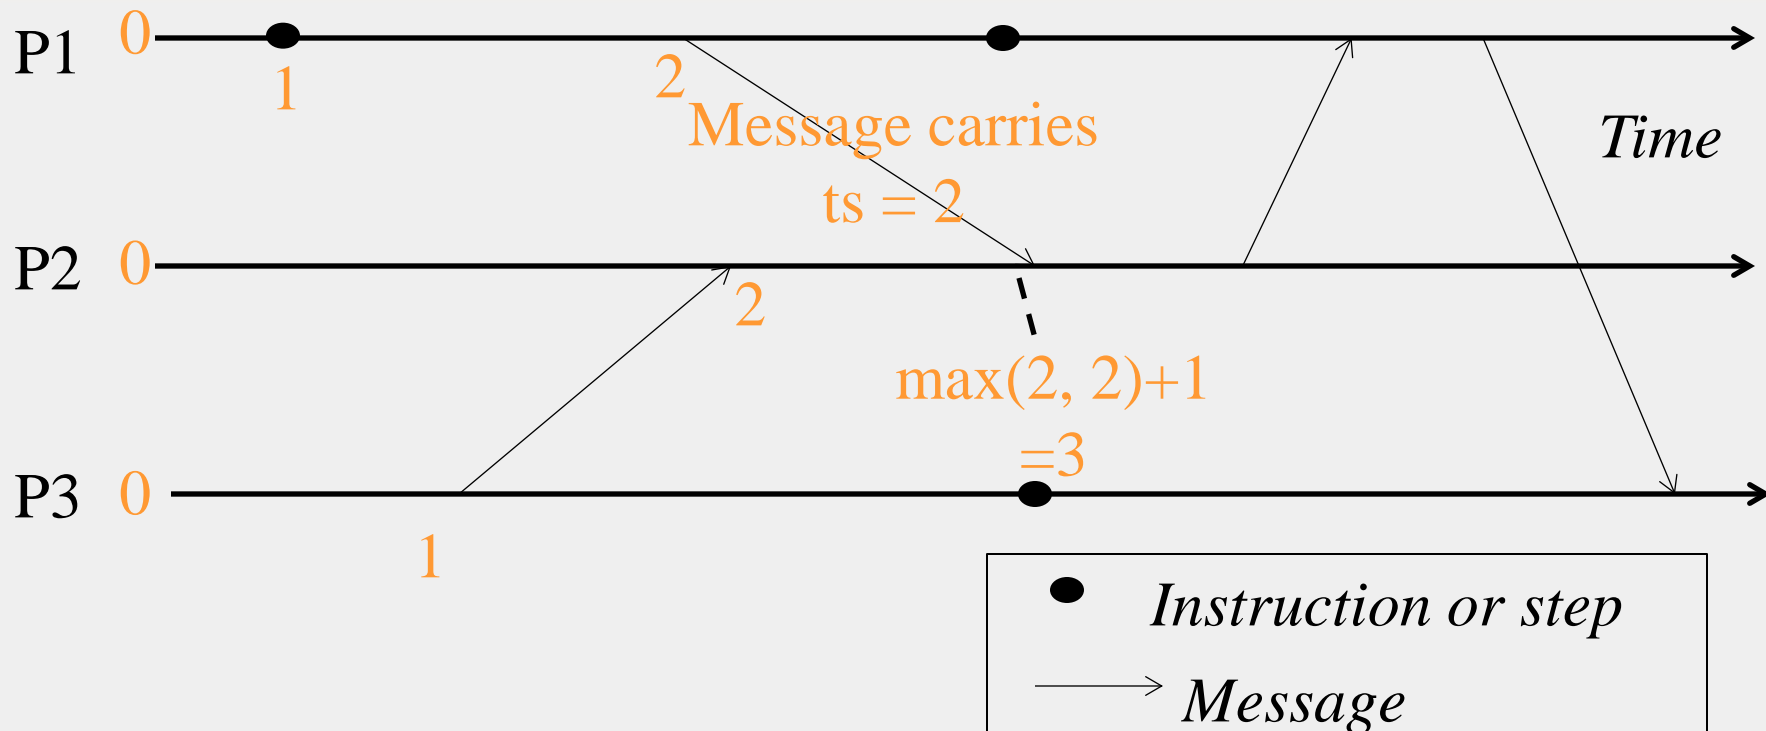
LAMPORT TIMESTAMPS



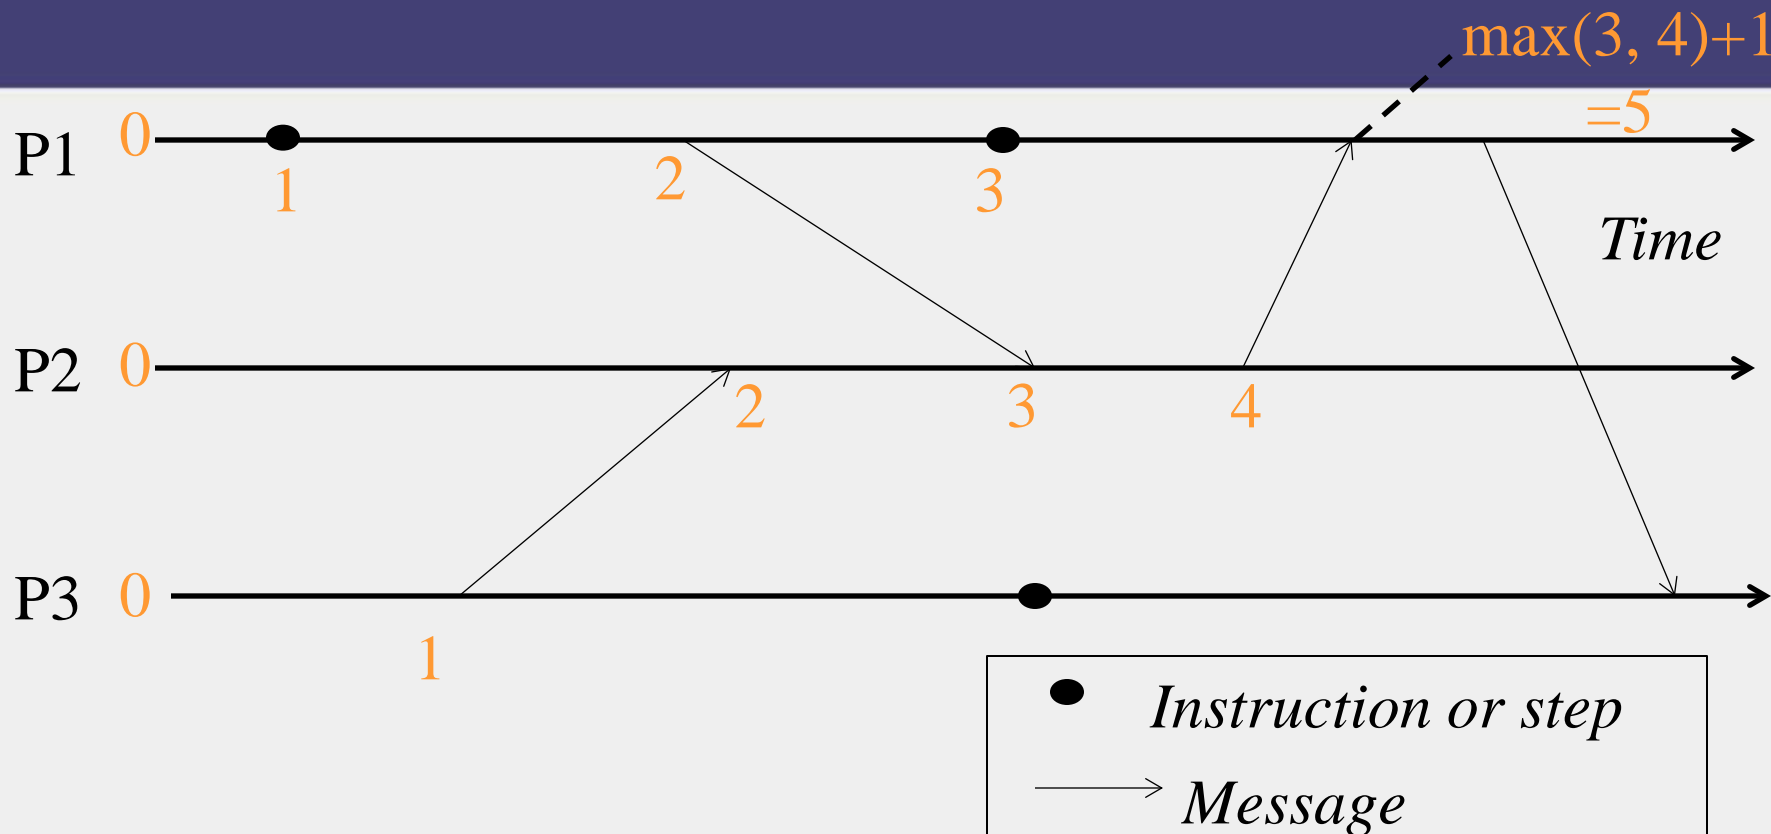
LAMPORT TIMESTAMPS



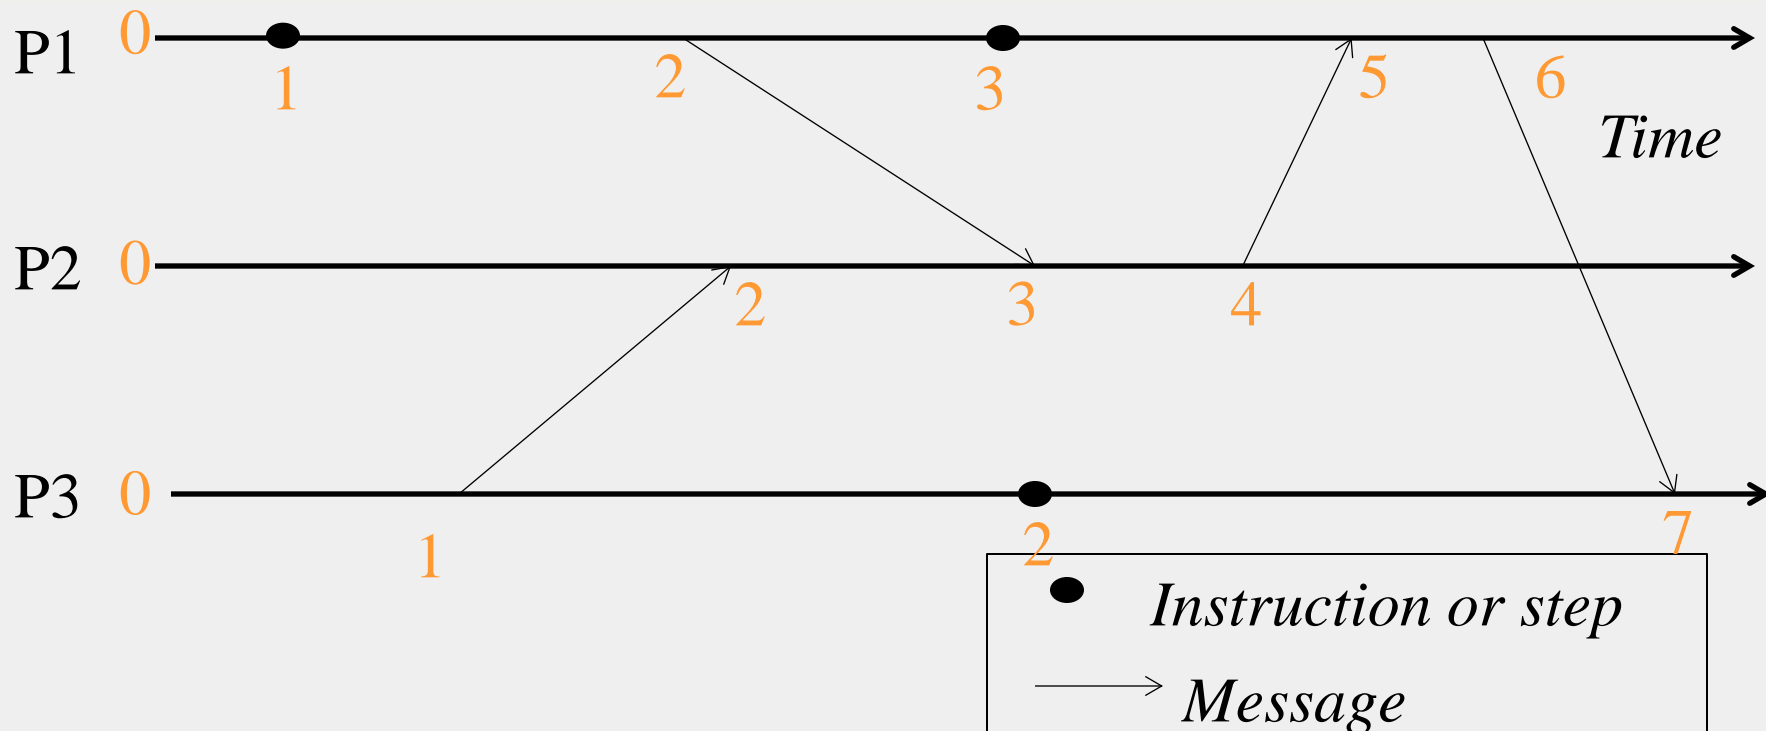
LAMPORT TIMESTAMPS




LAMPORT TIMESTAMPS



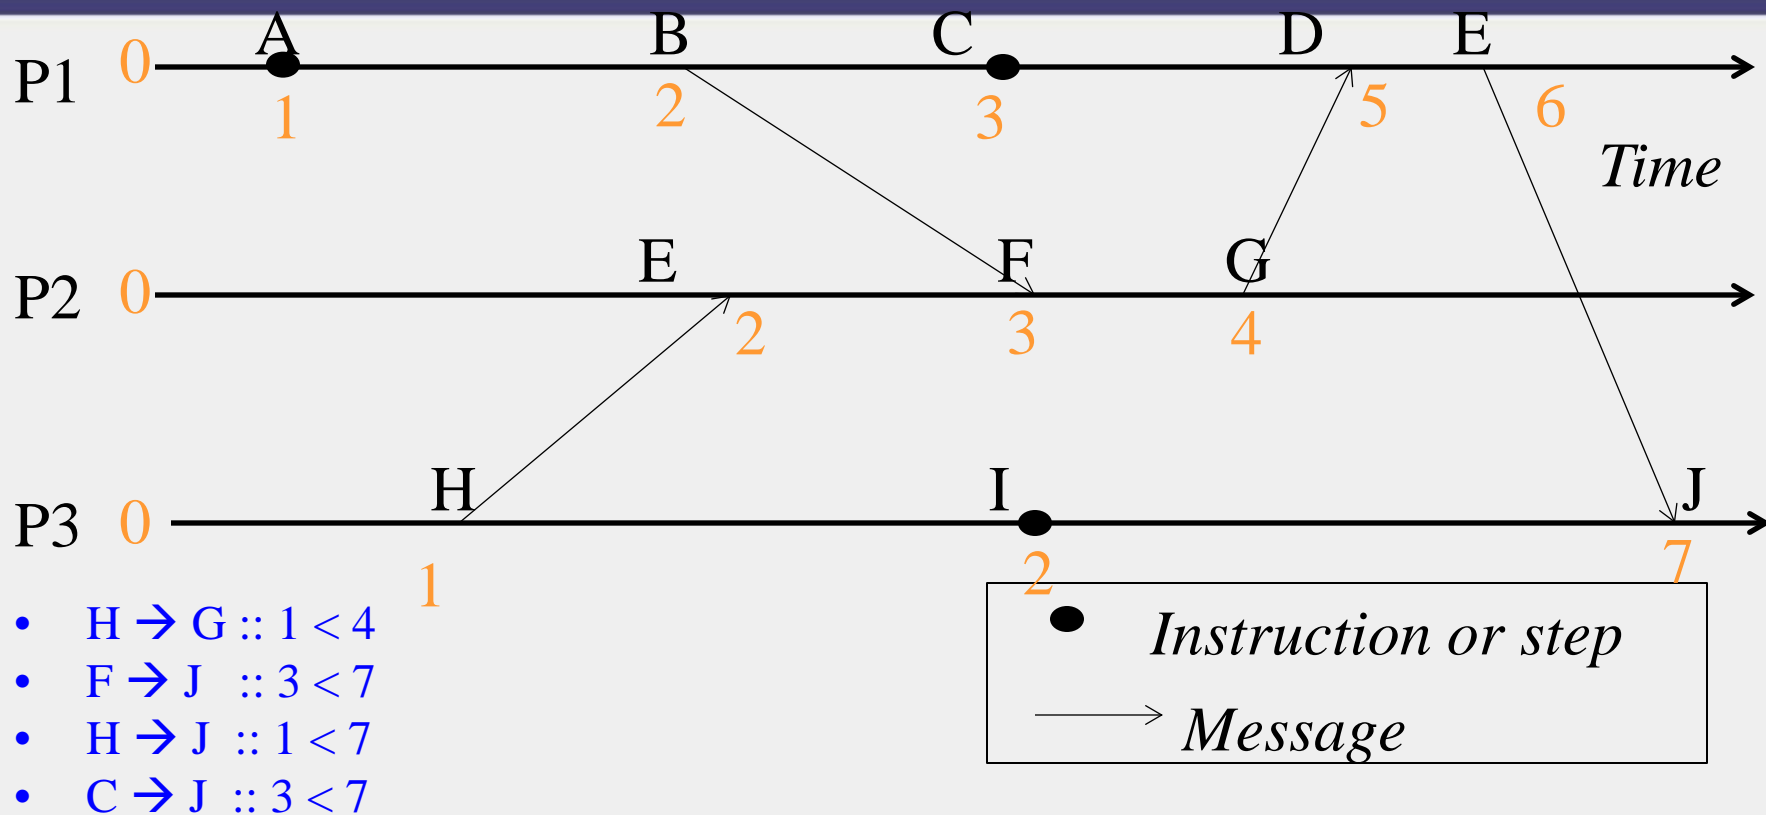
LAMPORT TIMESTAMPS




OBEYING CAUSALITY



OBEYING CAUSALITY (2)



NOT ALWAYS IMPLYING CAUSALITY



- $? C \rightarrow F ? :: 3 = 3$
- $? H \rightarrow C ? :: 1 < 3$
- (C, F) and (H, C) are pairs of concurrent events



CONCURRENT EVENTS

- A pair of concurrent events doesn't have a causal path from one event to another (either way, in the pair)
- Lamport timestamps not guaranteed to be ordered or unequal for concurrent events
- Ok, since concurrent events are not causality related!
- Remember

$E1 \rightarrow E2 \Rightarrow \text{timestamp}(E1) < \text{timestamp}(E2)$, BUT

$\text{timestamp}(E1) < \text{timestamp}(E2) \Rightarrow$

$\{E1 \rightarrow E2\} \text{ OR } \{E1 \text{ and } E2 \text{ concurrent}\}$

NEXT

- Can we have causal or logical timestamps from which we can tell if two events are concurrent or causally related?